1. Data preparation

Runs once (and when new reviews are imported). Everything runs on the PC: data, embeddings and LLM.

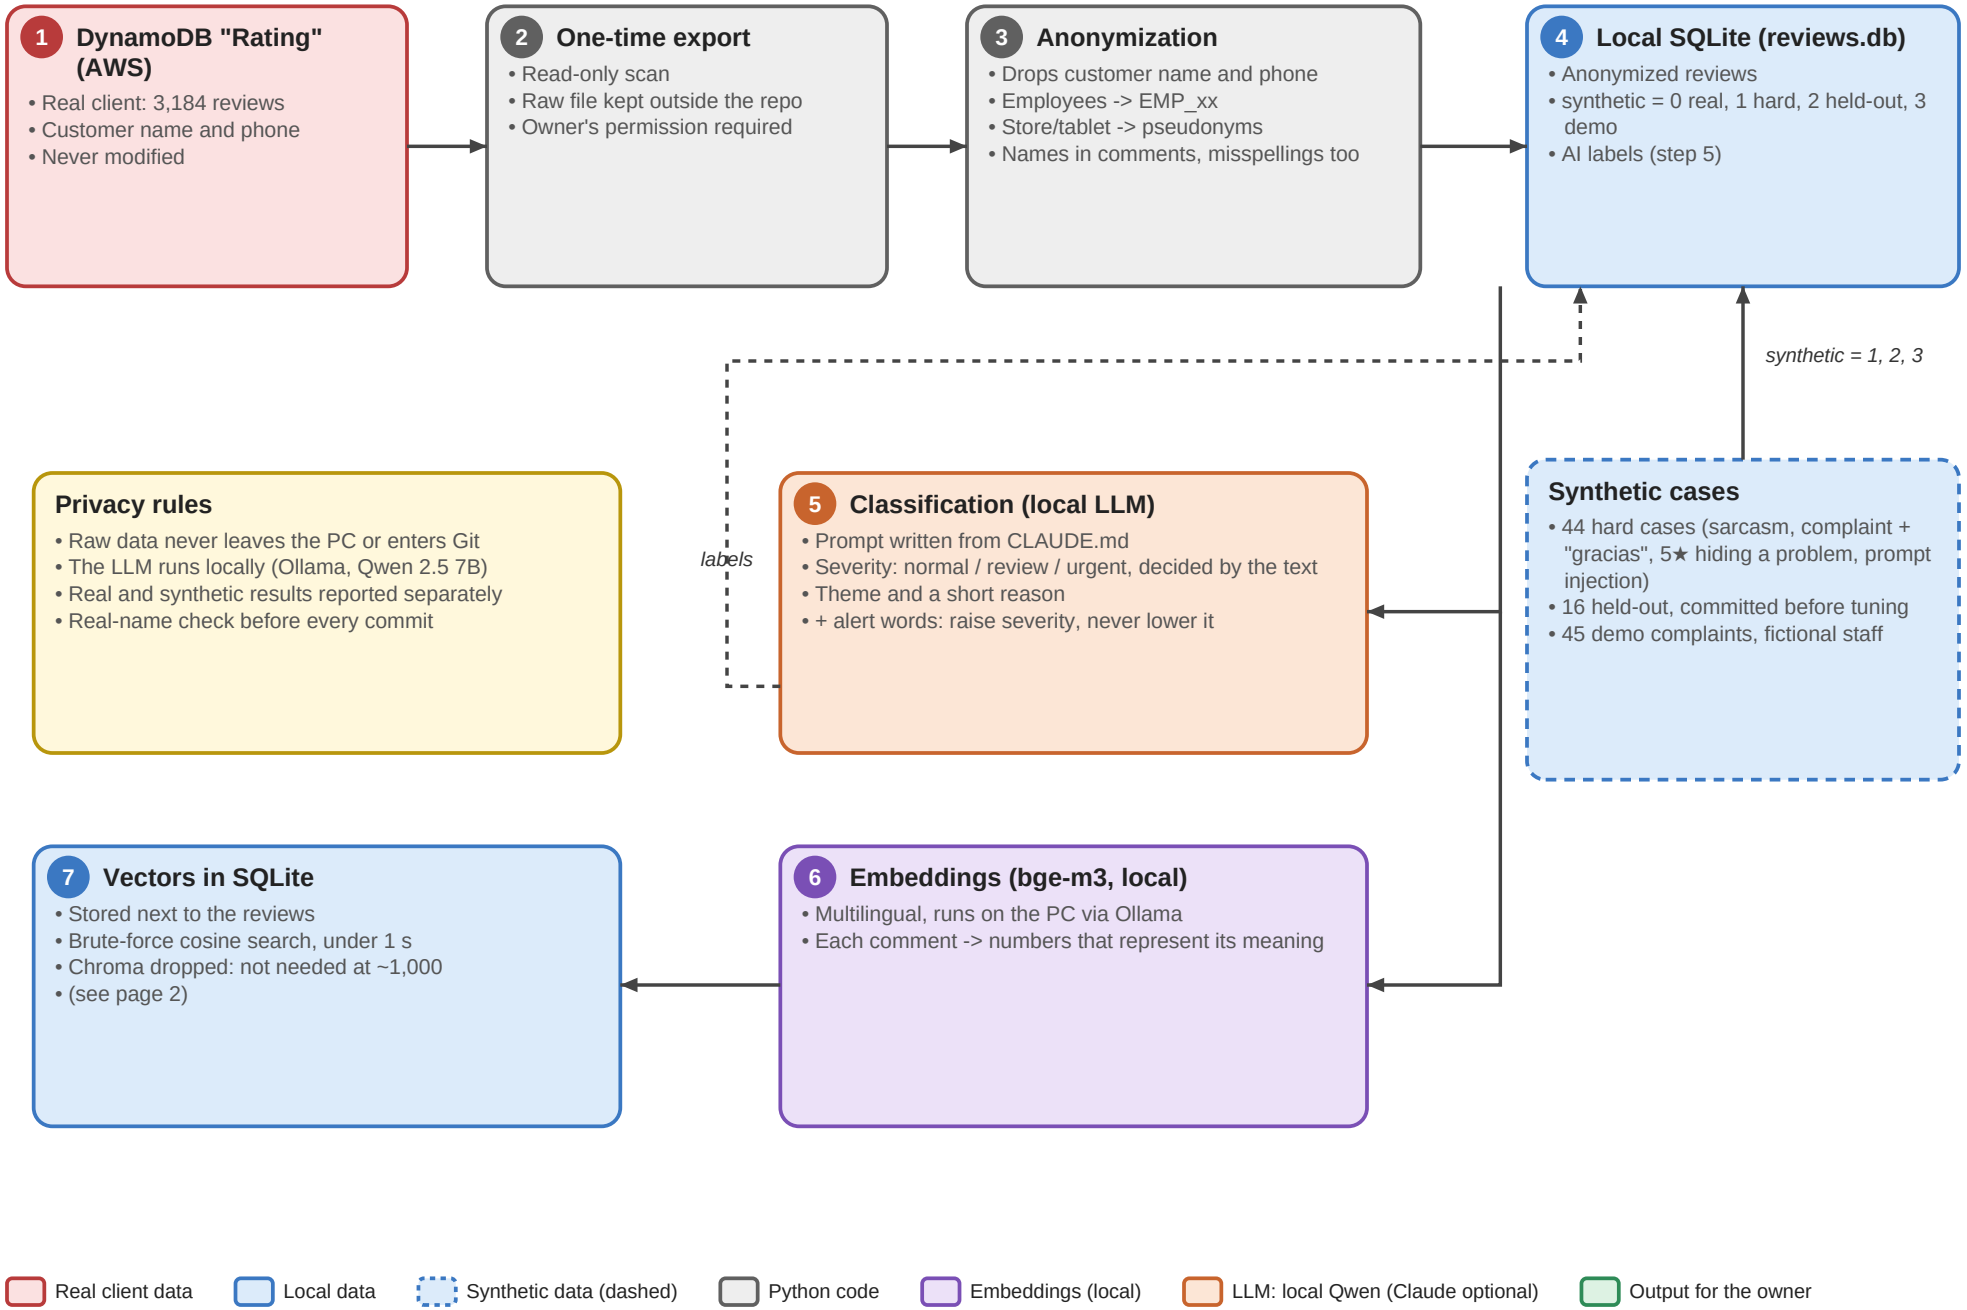

2. Owner questions ("ask your reviews")

Each question is routed: counts are answered by code; meaning questions use retrieval and the local LLM.

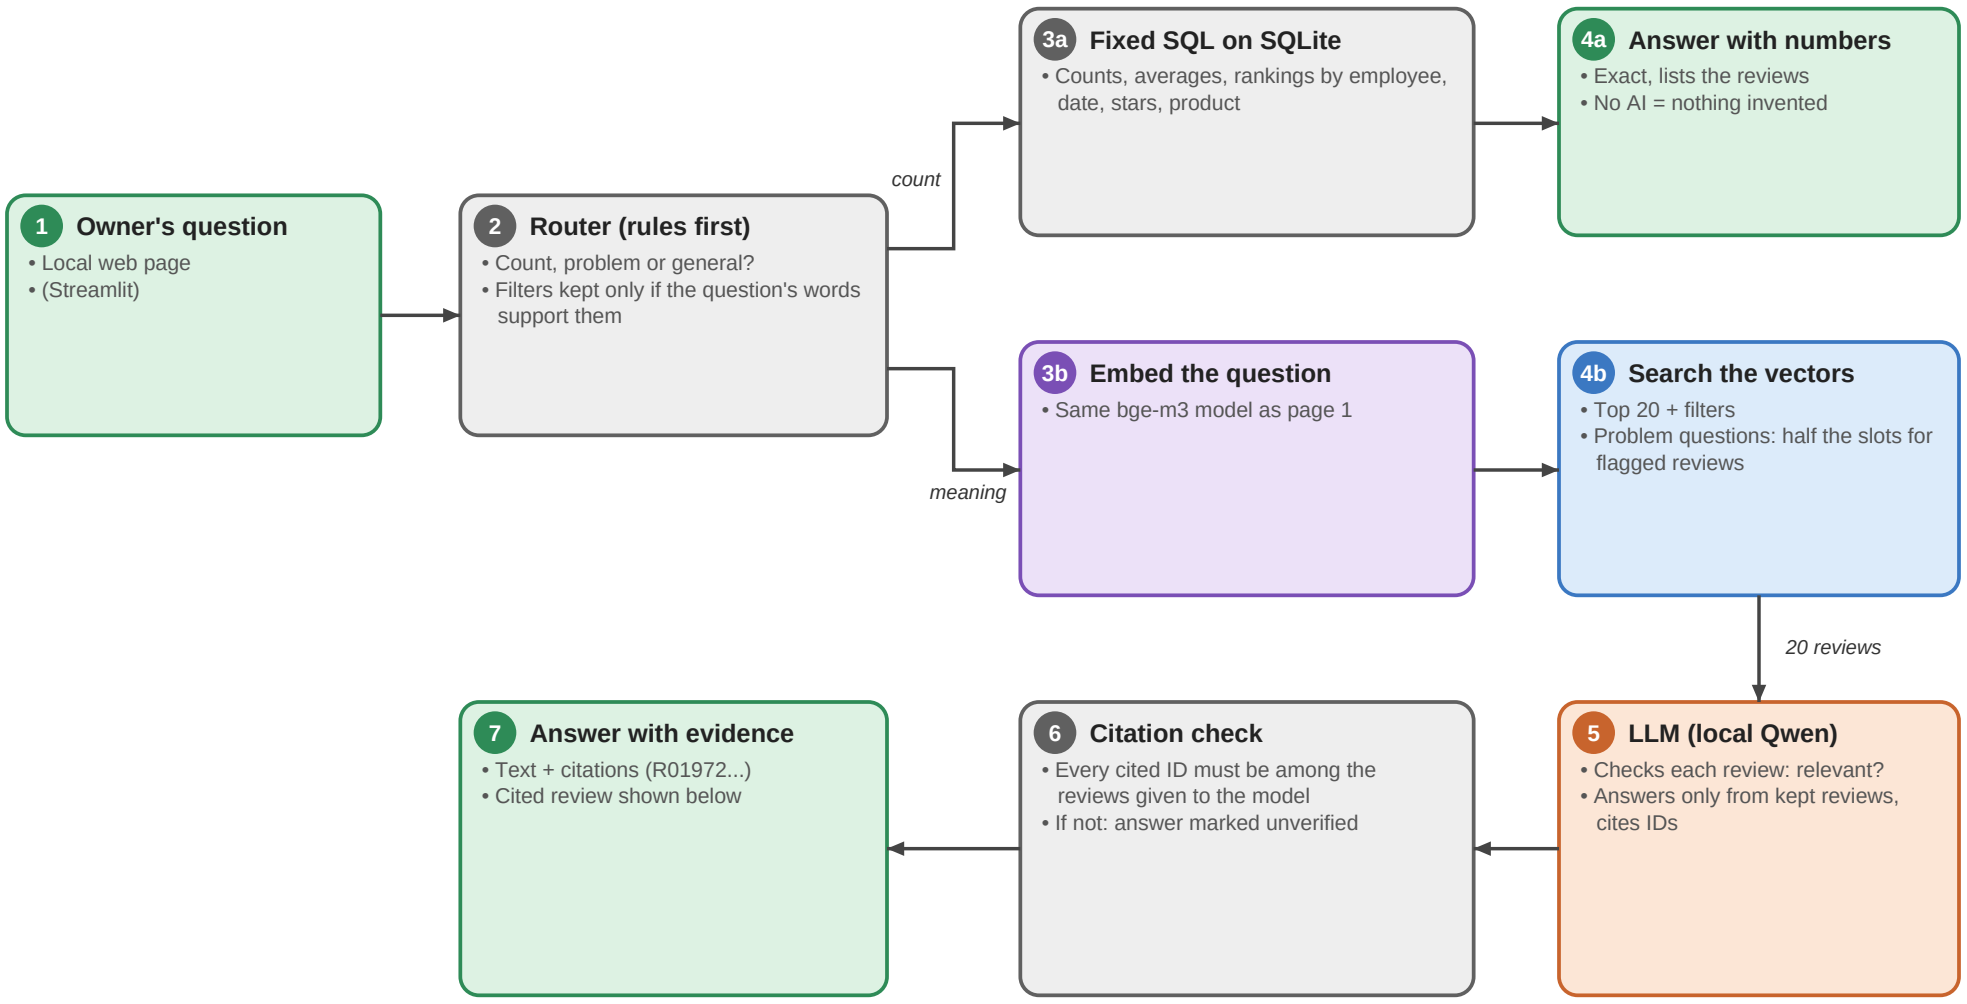

Examples (real output, docs/demo.md)

- "¿Cuántas reseñas de 1 estrella hubo en 2026?" -> count -> SQL -> 11, listing every review
- "¿Algún cliente tuvo un accidente o se cayó?" -> problem -> finds R01972 ("...despestañada y nos caimos") with no word in common
- Known failure: "¿Cuántas reseñas hubo en septiembre?" (no year) -> month filter dropped -> 3184 instead of 274

3. Evaluation, and the weekly digest that was cut

The evaluation measures whether the system works. The weekly digest was planned and not built.

Evaluation (evidence for the submission)

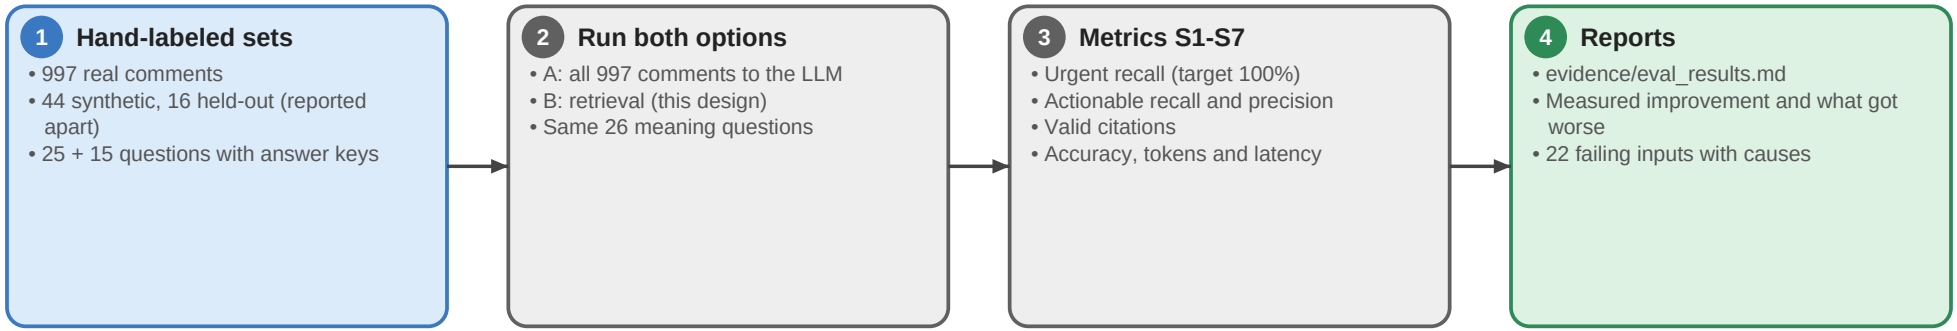

Weekly digest: planned, CUT (not built; see docs/effort.md)

Results against the definition of success (evidence/eval_results.md)

- Urgent recall: real 4/4, held-out 8/8, synthetic 15/17. Actionable: real 14/14 found, 14/14 precise (no false alarms)
- Counting questions: 13/13 exact within the sets; months without a year are silently dropped (F17)
- Meaning questions: 7/16 dev, 5/10 held-out; the 7B model drops reviews retrieval found
- 0 invalid citations in 40 answers. Full context (A) vs retrieval (B): 1/16 vs 7/16, 13 invented IDs vs 0